

AI Native Computing Stack and Ray Summit 2025Evolution of Computing eras:

Era	Compute Engine	Core stack/Technologies
Client server & Enterprise computing	Server (Windows/Unix)	Oracle databases etc
Internet and Web services	Virtual Machines (VMs)	LAMP stack (Linux Apache MySQL PHP)
Cloud computing	Container	Microservices, Big data, Kubernetes, MongoDB
AI Era	Ray (Distributed Computer Engine)	AI native computing stack (PyTorch, VLLM, Kubernetes)

Core Drivers demanding the AI Native Computing Stack:

- Heterogeneous hardware
 - CPUs + Accelerators
 - High speed networking
 - Purpose built AI data centres
- Complex, Heterogeneous AI processing Pipelines
 - Data processing
 - Training
 - Serving/Inference
- Non deterministic systems
- Higher iteration velocity

Ray: The AI Compute Engine

- Originated from Reinforcement Learning Research
- Adoption: 5 times more download increase since last year
- Core workloads:
 - Data processing
 - Training
 - Serving
- Evolution stages
 - Stage 1: Classic Neural Network (Moderate Scale)
 - Stage 2: First GenAI (Pre-Training Scaling Law, Accelerated Adoption)
 - Stage 3: Second GenAI (Post training - Reinforcement Learning, Massive Complexity)

AI Infrastructure Stack Layers:

- Top Layer: Training and Inference framework (PyTorch, VLM(Visual - Language model, integrates computer vision and natural language processing))
- Middle Layer: Distributed Compute Engine (Ray)
- Bottom Layer: Container orchestration (Kubernetes)
- Ecosystem: PyTorch Foundation

Challenge of Distributed Inference and Cross stack collaboration:

- Scale of inference
 - As models grow larger, running inference requires a lot more GPU
- Cross stack optimisation
 - Requires optimization across multiple stack levels, from API down to hardware placement
- Fine Grained Placement
 - Optimizing the placement of VLLM (Virtual Large Language Models)
- Cross node parallelism strategies
 - Requires establishing coordination across PyTorch, Ray, Kubernetes and VLLM

Ray | AnyscaleThe AI Compute Engine

- Supports any AI or ML workload
- Support any data types and model architectures
- Uses heterogenous GPUs and CPUs with fine grained, independent scaling
- Fully utilizes every accelerator
- Scale from your laptop to thousands of GPUs

Ray Features

- Parallel python code
- Multi modal data processing
- Model training
- Model serving
- Batch inference
- Reinforcement Learning
- Gen AI
- LLM inference
- LLM Fine Tuning

Ray integrations with pythonTasks

```
# Define the square task.
@ray.remote
def square(x):
    return x * x

# Launch four parallel square tasks.
futures = [square.remote(i) for i in range(4)]

# Retrieve results.
print(ray.get(futures))
# -> [0, 1, 4, 9]
```

Actors

- While tasks are stateless, Ray actors allows you to create stateful workers that maintain their internal state between method calls, when you instantiate a ray actor
 - Ray starts a dedicated worker process somewhere in your cluster
 - The Actors methods run on that specific worker and can access and modify its state
 - The actor executes method calls serially in the order it receives them, preserving consistency

```
# Define the Counter actor.
@ray.remote
class Counter:
    def __init__(self):
        self.i = 0

    def get(self):
        return self.i

    def incr(self, value):
        self.i += value

# Create a Counter actor.
c = Counter.remote()

# Submit calls to the actor. These calls run asynchronously but in
# submission order on the remote actor process.
for _ in range(10):
    c.incr.remote(1)

# Retrieve final actor state.
print(ray.get(c.get.remote()))
# -> 10
```

Objects

- Three main ways to work with objects in Ray
 - Implicit creation: When tasks and actors return values, they are automatically stored in ray's distributed object store
 - Explicit creation: Use ray.put() to directly place objects in the store
 - Passing references: Pass object references to other tasks and actors

```
import numpy as np

# Define a task that sums the values in a matrix.
@ray.remote
def sum_matrix(matrix):
    return np.sum(matrix)

# Call the task with a literal argument value.
print(ray.get(sum_matrix.remote(np.ones((100, 100)))))
# -> 10000.0

# Put a large array into the object store.
matrix_ref = ray.put(np.ones((1000, 1000)))

# Call the task with the object reference as an argument.
print(ray.get(sum_matrix.remote(matrix_ref)))
# -> 1000000.0
```

Ray Libraries

- Data: Scalable, framework agnostic data loading and transformation across training, tuning and prediction
- Train: Distributed multi node and multi core model training with fault tolerance that integrates with popular training libraries
- Tune: Scalable hyper parameter tuning to optimize model performance
- Serve: Scalable and programmable serving to deploy models for online interface

- RLLib: Scalable distributed reinforcement learning workloads

Anyscale Unified AI Platform:

- Goal: Productionization, Security, Governance
- Three main layers
 - Developer central
 - Workspaces
 - Debugging
 - Tooling
 - Ray Runtime (Optimized Ray Turbo Evolution)
 - Cluster controller (Compute Life Cycle Management)
- Platform innovations
 - Lineage Tracking
 - Anyscale Runtime (2x Faster Data, 7x Higher Throughput Serving)
 - GKE(Google Kubernetes Engine)/AKS(Azure Kubernetes Engine) integration
 - Multi Resource Clouds
 - Global Resource Scheduler: Flexible GPU allocation